

FINAL PRESENTATION · AI AGENTS

Job hunting, with an **agent** on your side

An AI career agent that finds the jobs, ranks them against your real experience, writes your CV and cover letter — and gets better every time you apply.

keinplankarriere.qantra.dev

THE PROBLEM

Applying for jobs is a **second job**

Fragmented

The same roles are scattered — and duplicated — across LinkedIn, StepStone, Xing and the Arbeitsagentur.

Repetitive

Every serious application means rewriting the CV and cover letter. 30–60 minutes, per job.


Blind

No feedback loop: which skills are in demand? Which of your applications actually worked?

Result: people either send generic applications everywhere — or burn out tailoring a few.

THE SOLUTION

One agent, the **whole loop**



Grounded in you

Upload your CV once — the AI turns it into a reviewed **experience base**. Every score and every document is built from what you've actually done. No fabrication, ever.

You stay in control

The agent prepares; **you** submit. No auto-apply, no stored job-board credentials, no ToS violations — a deliberate design choice.

FIND

Four boards, **one clean list**

LinkedIn

StepStone

Xing

Arbeitsagentur

- Search any query + location straight from the dashboard
- **Full job descriptions** fetched per posting — the fuel for everything downstream
- Skills auto-extracted (curated taxonomy, DE + EN)
- Fuzzy dedup merges the same role across boards

Trends, for free

Because everything is structured, the Trends tab shows what the market wants right now: top skills, salary ranges, remote split, hottest locations and companies — with a one-paragraph AI summary.

MATCH

Ranked for **you specifically**

Hybrid scoring

Layer 1 — a deterministic 0–100 score (skills 45 · role 20 · location · remote · seniority · salary): instant and free for every job.

Layer 2 — the top candidates get an LLM re-score that reads your actual experiences and *explains itself* in plain language.

It learns from outcomes

Mark a job as **interview** or **offer** and its skills become a success signal: similar jobs rise in the ranking, and the AI scorer is told what worked. Plus smart caching — unchanged jobs never pay for a second LLM call.

Every score comes with a visual breakdown — you always see *why* a job ranks where it does.

APPLY

From “interesting” to **application-ready** in one click

Tailored CV

your LaTeX résumé, every bullet rewritten for the job — nothing dropped

+

Cover letter

grounded in your experience, rendered to PDF

+

Apply kit

posting link + checklist + “mark as applied”

Make it yours

Your CV photo, your prompt add-ons (“emphasise leadership”, “max 250 words”), your own LaTeX template — all per user, all in Settings.

Always a valid PDF

Documents are compile-validated (Tectonic). If a tailored version fails, it falls back to your known-good template — generation never leaves you empty-handed.

Public. Multi-user. Try it right now.

[keinplankarriere](#)

Cloudflare Tunnel

Served from a small home desktop, exposed safely over HTTPS — no open ports, auto-restarting.

A workspace per visitor

Cookie sessions: your jobs, experiences, documents and settings are yours alone. Idle workspaces auto-clean after 48 h.

Bring your own AI

Visitors paste their own LLM key (Groq is free, OpenRouter, ...) — the host's credentials are never shared or exposed.

Fair-use queue

All heavy work runs through one serial worker with live “#2 in line” feedback — many users, one small machine, no meltdown.

<https://keinplankarriere.qantra.dev>

UNDER THE HOOD

Small stack, **built to survive**

The stack

FastAPI + SQLite · React + Vite · pure requests/BeautifulSoup scrapers · Tectonic for LaTeX→PDF · any OpenAI-compatible LLM (Kisski, Groq, OpenRouter, ...) · one Docker compose, two containers.

Resilience, learned the hard way

- **Self-healing LLM** — when the provider rotates models away, it re-picks an available one mid-request
- Deterministic core: dedup, skills and base scoring need no LLM at all
- Fallbacks everywhere: compile errors, rate limits, empty workspaces

Design principle: the LLM is reserved for judgment — parsing CVs, refining matches, writing documents. Everything else is plain, testable Python.

4 sprints · shipped & deployed

From zero to a public AI career agent

4

job boards, full descriptions

0–100

explainable match score

2 PDFs

CV + letter per job, one click

∞ users

isolated sessions + queue

Sprint 1: scrape & track → Sprint 2: match & profile → Sprint 3: generate & apply → Sprint 4: public, multi-user, self-improving.

keinplankarriere.qantra.dev – scan a job, get your CV